

ASH CLEMENTS

AI Red Team and Security Engineering

Phoenix, AZ | 480-604-3412 | ash@ashclements.dev | ashclements.dev | linkedin.com/in/ash-clements |
github.com/BadAsh99

Gray Swan Arena profile: app.grayswan.ai/arena/user/6a429ab1b2245b1643a7ffea

Open to hybrid, onsite San Francisco or Seattle about one week a month.

SUMMARY

I have spent 17 years in enterprise security, 9 of them at Palo Alto Networks, and for the past year I have also red-teamed production language models in Gray Swan's frontier-lab-funded arena, where placements pay and the top finishers are recruited for private engagements. What I keep finding is that the wins come from how an attack is structured, and almost never from how it is worded. Underneath the AI work is a full enterprise security career doing SASE and identity for Fortune 500 and regulated customers, plus live breach and zero-day response.

TECHNICAL SKILLS

AI Red Team: LLM red teaming, direct and indirect prompt injection, jailbreaks, multi-agent and RAG trust-boundary exploitation, multi-turn decomposition, adversarial evaluation, OWASP LLM Top 10, MITRE ATLAS

Offensive & AppSec: Red teaming, application security (static analysis, XXE, XSS, secret detection), threat modeling, vulnerability assessment, public-exploit validation and variant analysis, incident and breach response

Enterprise Security: SASE and Zero Trust (Prisma Access, ZTNA, SWG, CASB, DLP), identity (Cloud Identity Engine, User-ID), NGFW, PAN-OS, Panorama, identity-based policy enforcement

Development: Python, FastAPI, Flask, sentence-transformers, Docker, real-time SSE streaming, Next.js, TypeScript, Linux and bash (Ubuntu and other distributions, CLI-only)

Cloud & IaC: Azure, AWS, GCP, Terraform, multi-cloud misconfiguration and data-exposure assessment, CIS Benchmarks, OWASP Cloud Top 10

SIEM & SOAR: Log-forwarding architecture for customer SIEMs (syslog, CEF, LEEF, HTTPS and HEC), Strata Logging Service forwarding, Splunk, QRadar, Microsoft Sentinel, Exabeam, Cortex XSOAR

Certifications: PCNSE, 2024 · CISSP (In Progress, 2026)

AI SECURITY RESEARCH & RED-TEAM TOOLING

Independent research

- Gray Swan Arena, handle BadAsh99. I finished the Hazard Hunt challenge at #124 of 904 participants with 183 scored breaks, working the cyber track. Three of those were hard-model post-gate breaks, a category the platform accepted at a field rate of 0.32 percent. In the July indirect-injection challenge I finished #95 of 475 with 30 breaks.
- A single-shot request for a finished offensive artifact pattern-matches to malware and draws a categorical refusal, even from models that will write every component of it as a standalone benign function if asked over several turns. One short turn then wires together the functions the model already wrote. On the breaks that scored highest I supplied the premises and let the model author the step, which is what the judge credits as origination.
- I built breaks.jsonl out of my own runs to stop guessing at what actually worked. 128 records across 20 target behaviors, 51 of them individual scored breaks. Every row carries the gate type, about half carry the mechanism I used, and eight carry the judge's rationale word for word.
- AEGIS placed second in Protegrity's 2026 AI Pipeline Security Challenge. It is a RAG pipeline that tokenizes personal data on ingest and keeps it tokenized through embedding, the vector store and inference, so a successful injection walks away with tokens instead of people. Plaintext re-identified 100 percent of records on real MiniLM embeddings, AEGIS none, with identifier recall at 1.00.
- LLMGuardT2 fires 35 payloads across all 10 OWASP LLM Top 10 categories at live Claude and OpenAI endpoints, scoring each as vulnerable, partial, or resistant. Embedding similarity (all-MiniLM-L6-v2) catches paraphrased attacks a keyword scanner reads as clean.

- In badash-killchain, one prompt injection crosses agent and RAG trust boundaries inside a five-service microservices setup and reaches systems the attacker never touched directly. Mapped to MITRE ATLAS.
- AISeal covers the OWASP LLM Top 10 and the Agentic Top 10 (ASIO1-ASI10), pointing every finding at the EU AI Act article it breaks, somewhere in 9 through 15. CloudGuard scans AWS, Azure and GCP for misconfiguration and data exposure against CIS Benchmarks, read-only by default.

PROFESSIONAL EXPERIENCE

Palo Alto Networks | Remote, Phoenix, AZ | June 2017 - Present

Sr. Professional Services Consultant, SASE · 2021 - Present

Sr. Technical Support Engineer, SASE · 2017 - 2021

- Design and deploy Prisma Access and Zero Trust architectures for Fortune 500, energy and telecom customers, owning the identity layer on those engagements across Cloud Identity Engine, User-ID redistribution, DLP and policy enforcement.
- Validate public exploits and their variants almost daily, and support customer breach and zero-day incident response.
- Led a 50,000-device GlobalProtect and MDM rollout at a large telecom. Separately, took over an at-risk Fortune 500 deployment and saved a multi-million-dollar account.
- Built SCMRReady, internal tooling that runs automated pre-cutover checks on Panorama to Strata Cloud Manager migrations. Two customers deployed it in their own dev environments, a major software company and an online gambling operator, each after their own security team evaluated and approved it.
- Selected to lead a four-session internal enablement series on Panorama to Strata Cloud Manager migration for about 30 field engineers.
- Ran proof-of-value work on Prisma AIRS, the AI runtime security platform. Got the first live scan to a clean result in September 2026, after weeks of it failing on a profile-store mismatch nobody had documented.
- Own complex Prisma Access escalations, reproducing the hard ones in the lab and pushing the repro to engineering. Several have become shipped bug fixes, and I wrote the internal KBs and customer-facing troubleshooting guides that came out of them.

Network Security Engineer

American Express | Phoenix, AZ | Feb 2016 - May 2017

- Migrated enterprise data center firewalls from legacy Cisco ASA to Palo Alto NGFW using automated rule conversion and validation. Ran the HA clusters and kept them compliant around the clock, production and non-prod.

Earlier experience (2009 - 2016): network security and systems engineering roles covering vulnerability assessment, device hardening, and firewall operations.

EDUCATION

- Coursework in Information Security Technology, ECPI College of Technology, Richmond, VA
- Coursework, J. Sargeant Reynolds Community College, Richmond, VA